intel | aws

AWS powered by Intel. Co-engineered at the processor, firmware, Nitro, and hypervisor layers — and brought to market together: joint account teams, certified benchmarks per workload, funded proof-of-concept, and production references across every major enterprise solution area.

SAP & DATABASES | GENAI & INFERENCE | VMWARE & HYBRID | EDA · HPC · ANALYTICS

Solution description

- **Intel Xeon 6 instances (C8i, M8i, R8i) are built exclusively for AWS** — co-engineered silicon available in no other cloud or on-premises catalog.
- **2.5× memory bandwidth** versus prior generation compresses batch windows, reduces instance-hours, and accelerates memory-bound workloads — with no application changes.
- The same silicon extends on-premises: **AWS Outposts 2nd Gen ships Intel Xeon 6**, and Nutanix NC2 on AWS runs exclusively on Intel bare metal.

Challenges we solve

- **SAP, Oracle, and database estates** face renewal pressure and aging hardware — but migration fear is performance regression. Certified benchmarks remove the guesswork.
- **EDA, CFD, and HPC teams** are constrained by on-prem capacity and queue times — elastic burst on ParallelCluster eliminates the wait.
- **GenAI adoption stalls on GPU cost and availability** — Intel AMX runs RAG, vector DB, and enterprise assistants CPU-first, with predictable TCO.
- **VMware and OEM-displaced customers** need an exit that doesn't require refactoring — Intel-based AWS is the lift-and-shift path.

How we go to market together

- **Joint account teams:** Intel engineering specialists work alongside your AWS account team on sizing, architecture, and migration planning.
- **Benchmark before you commit:** free production benchmarking and TCO modeling on your actual workloads — plus up to 90 days of funded proof-of-concept compute.
- **ISV ecosystem in the deal:** SAP-certified sizing, EDA partnerships with Synopsys, Cadence, and Siemens EDA, HPC partnerships with Altair, Ansys, and Siemens Simcenter.
- **Intel runs on what it sells:** 1.2M EDA cores migrated to AWS, iGPT on Bedrock at 8.8B tokens monthly, and Intel's own SAP RISE migration as live references.

Customer wins

TOYOTA · AT&T · Disney+
BMW · J&J · JPMorgan Chase
CrowdStrike · Datadog · ServiceNow
JetZero · Hyundai Steel

Use cases — by workload

- **SAP modernization:** R8i at 142,100 aSAPS — 17% lower monthly cost, 27% faster throughput (Intel RISE migration)
- **Web & API tiers:** +60% NGINX throughput, 10–25% run-rate reduction (Disney+, AT&T)
- **Data & analytics:** batch windows compressed 2.5× (BMW, JPMorgan Chase, J&J — 35% faster genomics)
- **GenAI inference:** +52% SLM throughput, +45% tokens per dollar, no GPU (iGPT — 8.8B tokens/month)
- **EDA & HPC:** 1.2M Intel EDA cores on AWS; JetZero CFD, Hyundai Steel simulation on R8i
- **AI training at scale:** P5en — 35% lower collective comms latency (Sapphire Rapids + PCIe Gen5)

Every workload ships with a certified benchmark and a named production reference. Flip the page for the Intel Acceleration Program — funded migrations, proof-of-concept compute, and how to engage.

✓ AWS-Exclusive Intel Xeon 6
✓ 2.5× Memory Bandwidth
✓ Certified Benchmarks per Workload
✓ No Refactoring Required

FUNDED MIGRATIONS | 90-DAY FUNDED POC | FREE BENCHMARKING

Intel Acceleration Program (IAP): Funding for Every Motion

IAP puts Intel investment directly behind your migration — cash funding, proof-of-concept compute, certified benchmarks, and engineering support across every AWS service powered by Intel silicon: EC2, RDS, ElastiCache, EMR, ECS/EKS, Outposts, and P5en.

Funding is calculated on actual spend and disbursed to the partner or directly to the customer. Intel covers benchmark validation and TCO modeling on your workloads before any commitment.

THREE PROGRAMS — MATCHED TO WHERE YOU ARE TODAY

Program 1 · Migrate

4.5% of ARR

Net-new workloads and full migrations from on-premises or another cloud. Covers EC2, RDS, ElastiCache, EMR, EKS.

Program 1B · Upgrade

2.5% of ARR

Moving from 7th-gen Intel (c7i/m7i/r7i) to 8th-gen (c8i/m8i/r8i). 2× performance, 2.5× memory bandwidth. Instance family swap — no refactoring.

Program 2 · Optimize

1% of ARR

Right-size existing compute and free up budget. Savings fund the next workload — optimization becomes expansion.

1. **Qualify** — Identify the Intel-powered workload and the right program motion.
2. **Size** — Funding sized on actual spend — measured, not estimated.
3. **Engage** — Your AWS account team brings Intel into the opportunity. 4-day approval.
4. **SOW** — Statement of work issued. Funding disbursed to partner or customer.
5. **Validate** — Proof of performance confirms results. Funding completes.

STACK EVERY FUNDING VEHICLE ON THE SAME MIGRATION

Intel IAP		AWS MAP		VMware SPI		Greenfield SPI		PoC Funding
4.5 / 2.5 / 1%	+	Stackable	+	10%	+	10%	+	10% Y1

Registered Opportunity
Every funded engagement starts as a registered opportunity with your AWS account team.

Actual Spend
Funding is calculated on real, measured consumption — not list price estimates.

Fast Approval
Funding decisions in 4 days — aligned to your project timeline, not the other way around.

Proof of Performance
Results validated against the benchmark targets agreed at the start.

Intel-Powered Services
Eligible across Intel-powered EC2, RDS, ElastiCache, EMR, EKS, and Outposts.

Direct Disbursement
Funding goes to the partner delivering the work, or directly to the customer.

Intel covers up to 90 days of funded proof-of-concept compute, free production benchmarking, and TCO modeling — before you commit a dollar. Ask your AWS account team to bring Intel into the conversation.

✓ Funding Disbursed to Partner or Customer
✓ 4-Day Approval
✓ Stack IAP + MAP + SPI + PoC
✓ Benchmark Before You Commit