



AWS powered by Intel. Co-engineered at the processor, firmware, Nitro, and hypervisor layers — Intel Xeon 6 instances built exclusively for AWS, with certified benchmarks and production references across every major enterprise workload.

2.5×

MEMORY
BANDWIDTH VS
PRIOR GEN

+60%

NGINX
THROUGHPUT ON
C8I

+52%

SLM INFERENCE
THROUGHPUT

+45%

TOKENS PER
DOLLAR

142,100

SAP ASAPS —
HIGHEST AT LAUNCH

17%

LOWER SAP
MONTHLY COST ON
R8I

SAP Modernization

R8i and M8i deliver **17% lower monthly compute cost and 27% faster transaction throughput** versus 4th Gen Intel. R8i achieved **142,100 aSAPS** — highest SAP HANA certification at launch — with the broadest BTP integration on any cloud.

Live reference: Intel's own RISE migration, with SAP-certified benchmarks and sizing guidance.

VMware Transformation & Hybrid Cloud

Move from VMware or supply-constrained on-premises to Intel-based AWS **without refactoring**. Nutanix NC2 on AWS runs exclusively on Intel bare metal. AWS Outposts delivers **Intel Xeon 6 on-premises**.

Live reference: AT&T modernized IT and OSS/BSS on Intel Outposts SKUs. OEM displacement closes in weeks.

Generative AI & Enterprise Inference — CPU-First

C8i and R8i deliver **52% higher SLM inference throughput and 45% better tokens per dollar**. Intel AMX accelerates RAG, vector DB queries, and enterprise assistants — **without GPU dependency**.

Live reference: Intel's iGPT migrated to Bedrock in five weeks; now processing 8.8B tokens monthly.

EDA Modernization

R8i for memory-bound signoff, C8i for controllers, X2iezn for high-clock flows. **2.5× memory bandwidth reduces instance-hours per tape-out job**.

Live reference: Intel migrated 1.2M EDA cores to AWS — the largest EDA cloud transformation on record. Siemens EDA Calibre on R8i is the primary ISV reference.

HPC & Engineering Simulation

R8i for memory-intensive solves, C8i for meshing, X2iezn for high-clock niches. **ParallelCluster elastic burst eliminates queue time**.

Live references: JetZero runs blended-wing-body CFD with Altair on R8i targeting 50% fuel reduction. Hyundai Steel runs smart factory simulation on R8i and M8i.

Data & Analytics Modernization

R8i for ETL and analytics, C8i for orchestration. **2.5× memory bandwidth compresses batch windows** with no application changes.

Live references: BMW Cloud Data Hub, J&J genomics (35% faster analysis), JPMorgan Chase data mesh, CrowdStrike ML security.

Web, API & Cloud-Native Microservices

C8i delivers **60% higher NGINX throughput, driving 10–25% frontend run-rate reduction** on web and API tiers.

Live references: Disney+ runs realtime analytics and personalization on C8i and R8i. AT&T achieved 10–20% service-tier run-rate reduction.

GPU-Adjacent AI Training

P4d, P5, P5e, and P5en pair NVIDIA GPUs with Intel Xeon Scalable host CPUs. P5en delivers **35% lower collective communications latency** through Sapphire Rapids with PCIe Gen5 — reducing GPU starvation at scale.

Where it matters: data staging and preprocessing for distributed AI training at scale.



AWS-Exclusive Intel Xeon 6 Silicon



Certified Benchmarks per Workload



Production References Across Industries



No Refactoring Required

Intel Acceleration Program (IAP): How the Funding Works

IAP is the funding engine behind the co-sell — cash rebates, funded PoC compute, benchmarks, and demand generation across every AWS service powered by Intel silicon: EC2, RDS, ElastiCache, EMR, ECS/EKS, Outposts, and P5en.

Rebates are calculated on actual post-discount ARR, never list price — CUR verified. Cash is disbursed to the partner or the customer. AWS does not participate in payment, and AWS sellers carry zero quota impact.

THREE PROGRAMS — PICK THE MOTION THAT MATCHES THE CUSTOMER

Program 1 · Migrate

4.5% of ARR

Net-new workloads and full migrations from on-prem, GCP, or Azure. Covers EC2, RDS, ElastiCache, EMR, EKS.

Program 1B · Upgrade

2.5% of ARR

7th-gen Intel (c7i/m7i/r7i) to 8th-gen (c8i/m8i/r8i). 2× performance, 2.5× memory bandwidth. No refactoring.

Program 2 · Optimize

1% of ARR

Right-size and upsell pre-8th gen customers. Lock PPA — freed spend drives incremental workloads.

1

Qualify

Confirm Intel-eligible instance. ACE opportunity must exist first.

2

Size

Run calculator on actual post-discount spend. Cost Explorer baseline.

3

Engage

Submit via Deal Hub. Funded Head or PDM confirms. 4-day approval.

4

SOW

Generate SOW. Funding disbursed to partner or customer.

5

POP + Rebate

Proof of Performance by claim deadline. Rebate paid.

STACK EVERY FUNDING VEHICLE ON THE SAME DEAL

Intel IAP
4.5 / 2.5 / 1%

+

MAP
Stackable

+

VMware SPI
10%

+

Greenfield SPI
10%

+

PoC Funding
10% Y1

ACE First

Opportunity in ACE before funding. Cost Explorer baseline required.

Actual Spend

Rebate on post-discount ARR, never list price. CUR verified.

Deal Hub

All funding requests via Deal Hub. 4-day approval cycle.

POP Required

Proof of Performance before claim date or rolls to next cycle.

Intel Xeon Only

Intel-optimized instances. Trainium and Graviton excluded.

Zero Quota Impact

AWS sellers carry zero quota impact. Intel takes cloud goals.

Intel covers up to 90 days of funded PoC compute, free production benchmarking, and TCO modeling — before the customer commits a dollar. Bring Intel in before the ACE submission and let the funding do the convincing.



4-Day Deal Hub Approval



Stack IAP + MAP + SPI + PoC



Zero Quota Impact · AWS Sellers



Highspot: Intel DC Migration